

The Short-Time Fourier Transform

Xavier Serra

Music Technology Group

Universitat Pompeu Fabra, Barcelona

<http://mtg.upf.edu>

Index

- Short-time Fourier Transform equation
- Window type
- Window size
- FFT size
- Hop size
- Time-frequency compromise
- Inverse STFT
- STFT implementation

Short-time Fourier Transform


$$X_l[k] = \sum_{n=-N/2}^{N/2-1} w[n] x[n+lH] e^{-j2\pi kn/N} \quad l=0,1,\dots,$$

w: real window

l: number of frame

H: time advance of window (hop-size)




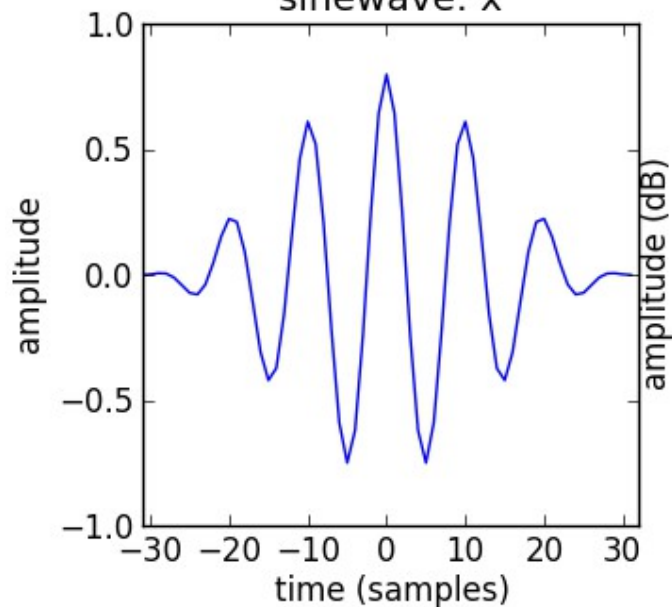


$$xw_l[n] = w[n]x[n+lH] \quad l=0,1,\dots,$$



Transform of a windowed sinewave

$$x[n] = A_0 \cos(2\pi k_0 n/N) = \frac{A_0}{2} e^{j2\pi k_0 n/N} + \frac{A_0}{2} e^{-j2\pi k_0 n/N}$$

$$\begin{aligned} X[k] &= \sum_{n=-N/2}^{N/2-1} w[n] x[n] e^{-j2\pi kn/N} \\ &= \sum_{n=-N/2}^{N/2-1} w[n] \left(\frac{A_0}{2} e^{j2\pi k_0 n/N} + \frac{A_0}{2} e^{-j2\pi k_0 n/N} \right) e^{-j2\pi kn/N} \\ &= \sum_{n=-N/2}^{N/2-1} w[n] \frac{A_0}{2} e^{j2\pi k_0 n/N} e^{-j2\pi kn/N} + \sum_{n=-N/2}^{N/2-1} w[n] \frac{A_0}{2} e^{-j2\pi k_0 n/N} e^{-j2\pi kn/N} \\ &= \frac{A_0}{2} \sum_{n=-N/2}^{N/2-1} w[n] e^{-j2\pi(k-k_0)n/N} + \frac{A_0}{2} \sum_{n=-N/2}^{N/2-1} w[n] e^{-j2\pi(k+k_0)n/N} \\ &= \frac{A_0}{2} W[k-k_0] + \frac{A_0}{2} W[k+k_0] \end{aligned}$$

hanning window: w magnitude spectrum: $\text{abs}(W)$ phase spectrum: $\text{angle}(W)$ sinewave: x magnitude spectrum: $\text{abs}(X)$ phase spectrum: $\text{angle}(X)$ 

```

N = 256
M = 63
f0 = 1000
fs = 10000
A0 = .8
hN = N/2
hM = (M+1)/2
x = A0 * np.cos(2*np.pi*f0/fs*np.arange(-hM+1,hM))

w = np.hanning(M)
plt.subplot(2,3,1)
plt.plot(np.arange(-hM+1, hM), w, 'b')

fftbuffer[:hM] = w[hM-1:]
fftbuffer[N-hM+1:] = w[:hM-1]
X = fft(fftbuffer)
X1[:hN] = X[hN:]
X1[N-hN:] = X[:hN]
mX = 20*np.log10(abs(X1))
plt.subplot(2,3,2)
plt.plot(np.arange(-hN, hN), mX, 'r')

pX = np.angle(X1)
plt.subplot(2,3,3)
plt.plot(np.arange(-hN, hN), np.unwrap(pX), 'g')

plt.subplot(2,3,4)
xw = x*w
plt.plot(np.arange(-hM+1, hM), xw, 'b')


fftbuffer[0:hM] = xw[hM-1:]
fftbuffer[N-hM+1:] = xw[:hM-1]
X = fft(fftbuffer)
X2[:hN] = X[hN:]
X2[N-hN:] = X[:hN]
mX2 = 20*np.log10(abs(X2))
plt.subplot(2,3,5)
plt.plot(np.arange(-hN, hN), mX2, 'r')

pX = np.angle(X2)
plt.subplot(2,3,6)
plt.plot(np.arange(-hN, hN), np.unwrap(pX), 'g')


```

Window type

All standard windows are real and symmetric and have a frequency spectrum with a sinc-like shape.



rectangular
window



magnitude
spectrum

The choice is mainly determined by two of the spectrum's characteristics:

1. Width of main lobe
2. Highest side-lobe level


Window functions in Scipy

<code>barthann (M[, sym])</code>	Return a modified Bartlett-Hann window.
<code>bartlett (M[, sym])</code>	Return a Bartlett window.
<code>blackman (M[, sym])</code>	Return a Blackman window.
<code>blackmanharris (M[, sym])</code>	Return a minimum 4-term Blackman-Harris window.
<code>bohman (M[, sym])</code>	Return a Bohman window.
<code>boxcar (M[, sym])</code>	Return a boxcar or rectangular window.
<code>chebwin (M, at[, sym])</code>	Return a Dolph-Chebyshev window.
<code>flattop (M[, sym])</code>	Return a flat top window.
<code>gaussian (M, std[, sym])</code>	Return a Gaussian window.
<code>general-gaussian (M, p, sig[, sym])</code>	Return a window with a generalized Gaussian shape.
<code>hamming (M[, sym])</code>	Return a Hamming window.
<code>hann (M[, sym])</code>	Return a Hann window.
<code>kaiser (M, beta[, sym])</code>	Return a Kaiser window.
<code>nuttall (M[, sym])</code>	Return a minimum 4-term Blackman-Harris window according to Nuttall.
<code>parzen (M[, sym])</code>	Return a Parzen window.
<code>slepian (M, width[, sym])</code>	Return a digital Slepian window.
<code>triang (M[, sym])</code>	Return a triangular window.

Rectangular window

$$w(n) = 1, \quad n = -M/2, \dots, 0, \dots, M/2$$
$$= 0, \quad \text{elsewhere}$$

$$W(\omega) = \frac{\sin(\omega M/2)}{\sin(\omega/2)}$$



main-lobe width: 2 bins
side-lobe level: -13.3 dB

```
import matplotlib.pyplot as plt
import numpy as np
from scipy.fftpack import fft

M = 64
N = 1024
hN = N/2
hM = M/2
fftbuffer = np.zeros(N)

fftbuffer[hN-hM:hN+hM]=np.ones(M)
plt.subplot(2,1,1)
plt.plot(np.arange(-hN, hN), fftbuffer, 'b')


X = fft(fftbuffer)
mX = 20*np.log10(abs(X))
mX1[:hN] = mX[hN:]
mX1[N-hN:] = mX[:hN]
plt.subplot(2,1,2)
plt.plot(np.arange(-hN, hN), mX1-max(mX), 'r')
```

Hanning window

$$w(n) = .5 + .5 \cos(2n\pi/M),$$
$$n = -M/2, \dots, 0, \dots, M/2$$

$$W(\omega) = .5D(\omega) +$$
$$.25 \left[D(\omega - 2\pi/N) + D(\omega + 2\pi/N) \right]$$

where $D(\omega) = \frac{\sin(\omega M/2)}{\sin(\omega/2)}$



main-lobe width: 4 bins
side-lobe level: -31.5 dB


Hamming window

$$w(n) = .54 + .46 \cos(2n\pi/M),$$

$$n = -M/2, \dots, 0, \dots, M/2$$

$$W(\omega) = .5D(\omega) +$$


$$.25 \left[D(\omega - 2\pi/N) + D(\omega + 2\pi/N) \right]$$



main-lobe width: 4 bins
side-lobe level: -42.7 dB

Blackman window

$$w[n] = 0.42 - 0.5 \cos(2\pi n/M) + 0.08 \cos(4\pi n/M)$$




main-lobe width: 6 bins
side-lobe level: -58 dB

Blackman-Harris window

minimum 4-term Blackman-Harris window

$$w(n) = \frac{1}{M} \sum_{l=0}^3 \alpha_l \cos(2nl\pi/M), \quad n = -M/2, \dots, 0, \dots, M/2$$

where $\alpha_0 = 0.35875, \alpha_1 = 0.48829, \alpha_2 = 0.14128, \alpha_3 = 0.01168$






main lobe width : 9 bins
side-lobe level : -92dB




Window size

Affects the time versus frequency resolution

$M = 128$



$M = 1024$




```
(fs, x) = read('oboe.wav')
N = 128
start = .8*fs
xw = x[start:start+N] * np.hamming(N)
plt.subplot(321)
plt.plot(np.arange(start, (start+N), 1.0)/fs, xw)

X = fft(xw)
mX = 20 * np.log10(abs(X)/N)
pX = np.unwrap(np.angle(X))
plt.subplot(323)
plt.plot(np.arange(0, fs/2.0, float(fs)/N), mX[0:N/2])





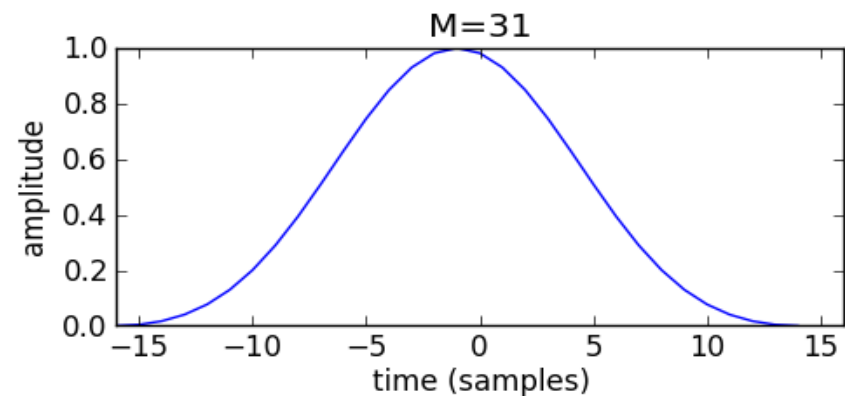

plt.subplot(325)
plt.plot(np.arange(0, fs/2.0, float(fs)/N), pX[:N/2])

N = 1024
start = .8*fs
xw = x[start:start+N] * np.hamming(N)
plt.subplot(322)
plt.plot(np.arange(start, (start+N), 1.0)/fs, xw)

X = fft(xw)
mX = 20 * np.log10(abs(X)/N)
pX = np.unwrap(np.angle(X))
plt.subplot(324)
plt.plot(np.arange(0, fs/2.0, float(fs)/N), mX[0:N/2])


plt.subplot(326)
plt.plot(np.arange(0, fs/2.0, float(fs)/N), pX[:N/2])
```

Even-Odd size Window




FFT size

FFT size, N , should be equal or larger than the window size, M




$M = 512$



$M = 512$

$N = 512$



$M = 512$

$N = 1024$

```
(fs, x) = read('oboe.wav')
M = 512
N = 512
start = .8*fs
xw = x[start:start+M] * np.hamming(M)
plt.subplot(311)
plt.plot(np.arange(start, (start+M), 1.0)/fs, xw)

X = fft(xw, N)
mX = 20 * np.log10(abs(X)/N)
pX = np.unwrap(np.angle(X))
plt.subplot(312)
plt.plot(np.arange(0, fs/2.0, float(fs)/N), mX[0:N/2])

M = 512
N = 1024
start = .8*fs
xw = x[start:start+M] * np.hamming(M)
X = fft(xw, N)
mX = 20 * np.log10(abs(X)/N)
plt.subplot(313)
plt.plot(np.arange(0, fs/2.0, float(fs)/N), mX[0:N/2])
```

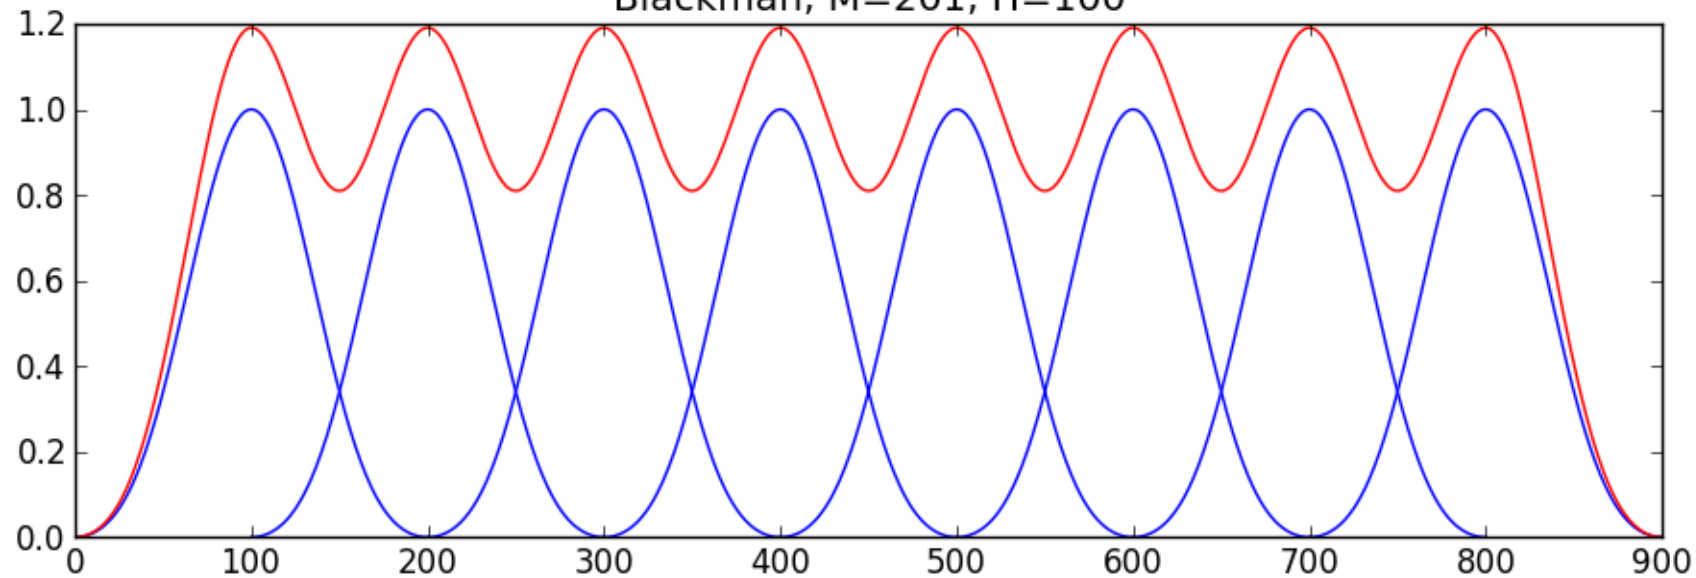
Hop size

- Successive frames should overlap in time in such a way that all the data are weighted equally.
- For certain windows exists perfect overlap factors. Rectangular: M/j , Hanning and Hamming: $(M/2)/j$, where $j = 0, 1, \dots$


The overlap factor can be expressed by:

$$A_w[n] = \sum_{l=0}^{L-1} w[n - lH] = c$$

Blackman, $M=201$, $H=100$



Blackman, $M=201$, $H=50$




```
N= 1000
M = 201
w = signal.blackman(M)
w1 = w/sum(w)

H = 100
pin = 0
pend = N - M
plt.subplot(211)
while pin<pend:
    y [pin:pin+M] += w1*w
    plt.plot(np.arange(pin, pin+M), w, 'b')
    pin += H
plt.plot(np.arange(0, N), y, 'r')


H = 50
pin = 0
pend = N - M
plt.subplot(212)
while pin<pend:
    y [pin:pin+M] += w1*w
    plt.plot(np.arange(pin, pin+M), w, 'b')
    pin += H
plt.plot(np.arange(0, N), y, 'r')
```

Time-frequency compromise

$N = 256$, window = hamming(256), $H = 128$



$N = 1024$, window = hamming(1024), $H = 128$




```
from pylab import *
from scipy.io.wavfile import read

(fs, x) = read('piano.wav')
subplot(2,1,1)
specgram(x,NFFT=256,window=hamming(256),Fs=fs,noverlap=128)

subplot(2,1,2)
specgram(x,NFFT=1024,window=hamming(1024),Fs=fs,noverlap=128)
```

Inverse STFT

$$y[n] = \sum_{l=0}^{L-1} \text{Shift}_{lH, n} \left[\frac{1}{N} \sum_{k=-N/2}^{N/2-1} X_l[k] e^{j2\pi kn/N} \right]$$


each synthesized frame is:

$$yw_l[n] = x(n + lH) w[n]$$

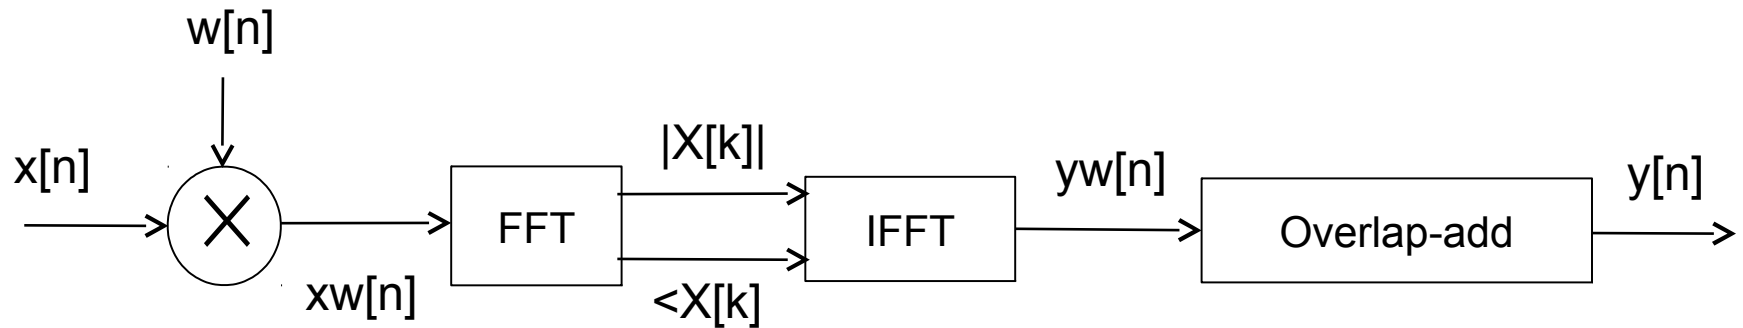
and the synthesized sound is:

$$y[n] = \sum_{l=0}^{L-1} yw_l[n] = x[n] \sum_{l=0}^{L-1} w[n - lH]$$

$$yw_l[n] = w[n] y[n + lH] \quad l = 0, 1, \dots,$$



STFT implementation diagram



```

def stft(x, fs, w, N, H) :
    hN = N/2
    hM1 = int(math.floor((w.size+1)/2))
    hM2 = int(math.floor(w.size/2))
    pin = hM1
    pend = x.size-hM1
    w = w / sum(w)
    while pin<pend:
        #-----analysis-----
        xw = x[pin-hM1:pin+hM2]*w
        fftbuffer[:hM1] = xw[hM2:]
        fftbuffer[N-hM2:] = xw[:hM2]
        X = fft(fftbuffer)
        mX = 20 * np.log10( abs(X[:hN]) )
        pX = np.unwrap( np.angle(X[:hN]) )
        #----synthesis----
        Y[:hN] = 10**(mX/20) * np.exp(1j*pX)
        Y[hN+1:]=10**(mX[:0:-1]/20)*np.exp(-1j*pX[:0:-1])
        fftbuffer = np.real( ifft(Y) )
        yw[:hM2] = fftbuffer[N-hM2:]
        yw[hM2:] = fftbuffer[:hM1]
        y[pin-hM1:pin+hM2] += H*yw
        pin += H
    return y

```

References

- <https://ccrma.stanford.edu/~jos/sasp/>
- <https://en.wikipedia.org/wiki/STFT>
- https://en.wikipedia.org/wiki/Window_function
- Full code of plots and accompanying labs available at:
<https://github.com/MTG/sms-tools>

Credits

All the slides of this presentation are released under an Attribution-Noncommercial-Share Alike license.